

Pathfinding

In Clojure



What is pathfinding?

- Find the shortest¹ path between two points
- More generally known as the shortest path problem
- Typically applied to graphs
- Can be applied to abstract problems
 - eg, finding a “path” through a graph of states to reach a desired state

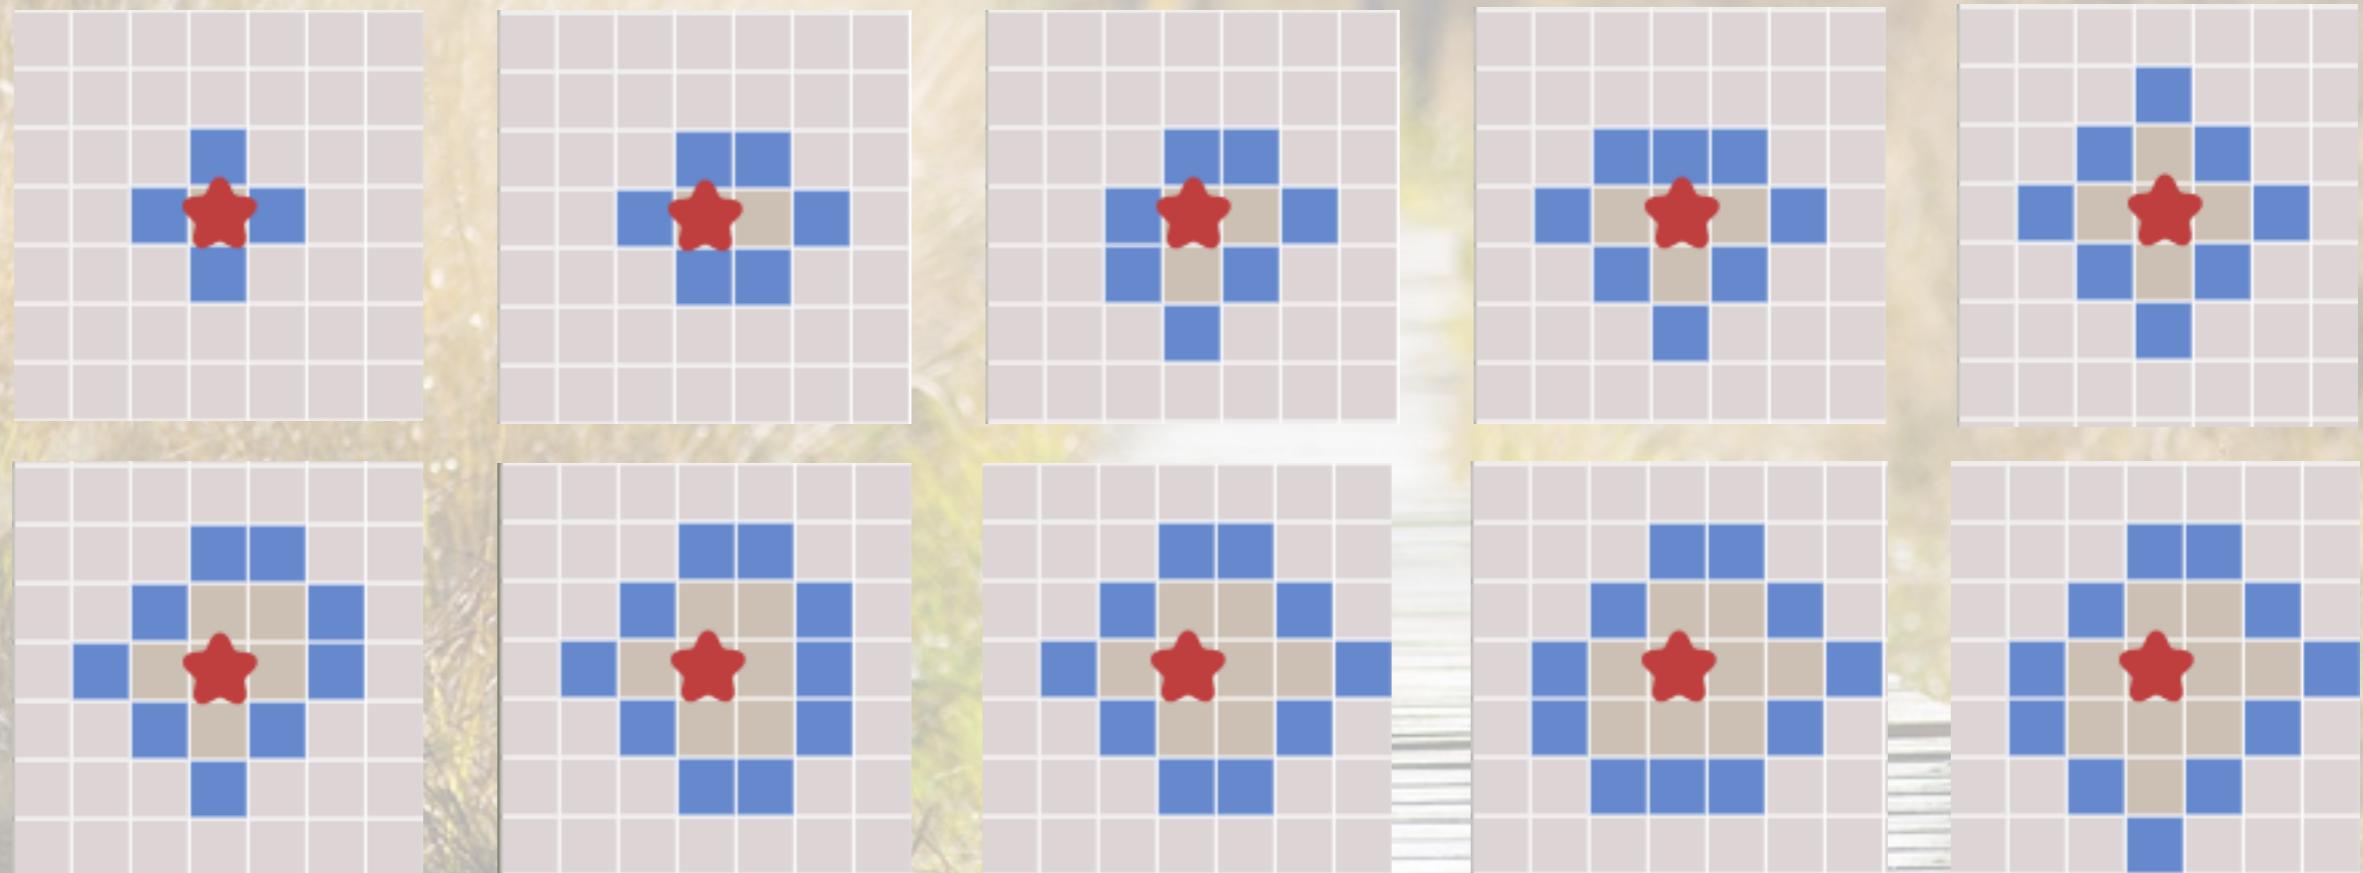
¹. Shortest or “cheapest”, for example, preferring large roads over small

Algorithms

- Breadth-first Search
- Dijkstra's Shortest Path
- A*
- Very similar algorithms, build on each other

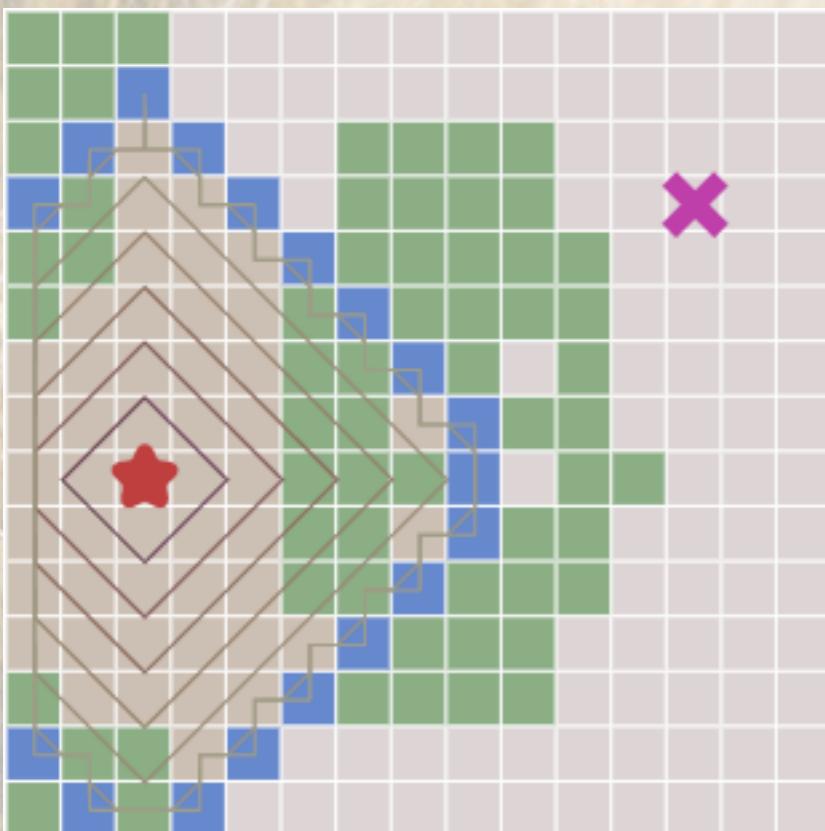
Breadth-first Search

- Explores each neighbouring “node” before exploring deeper

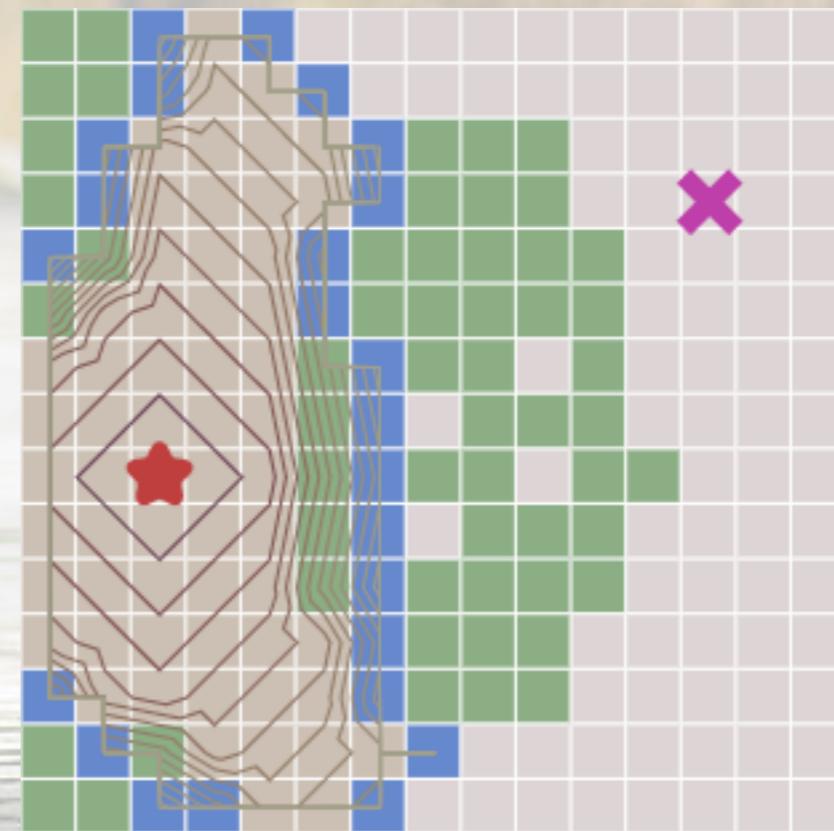


Dijkstra's Shortest Path

- Similar to BFS, but explores “cheapest” neighbours first



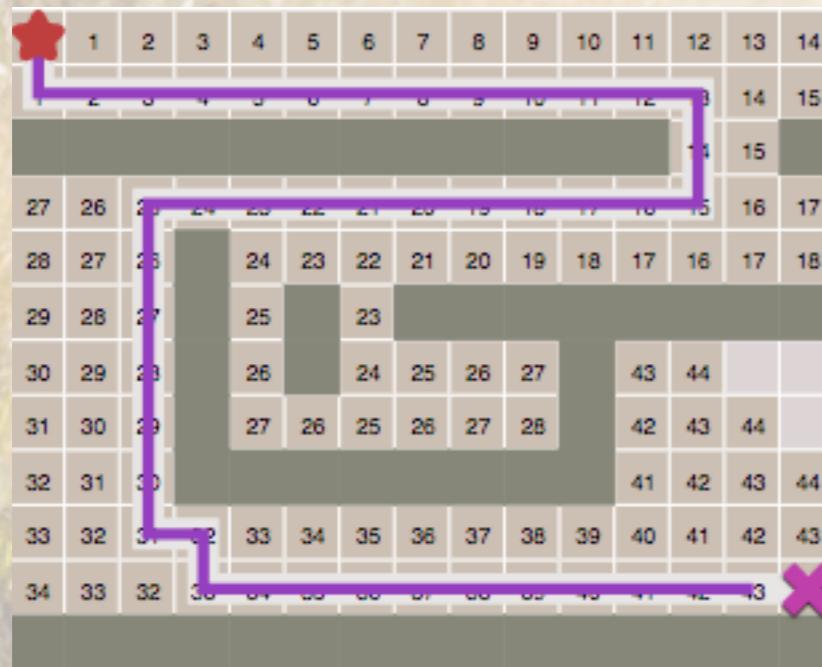
Breadth First



Dijkstra's algorithm

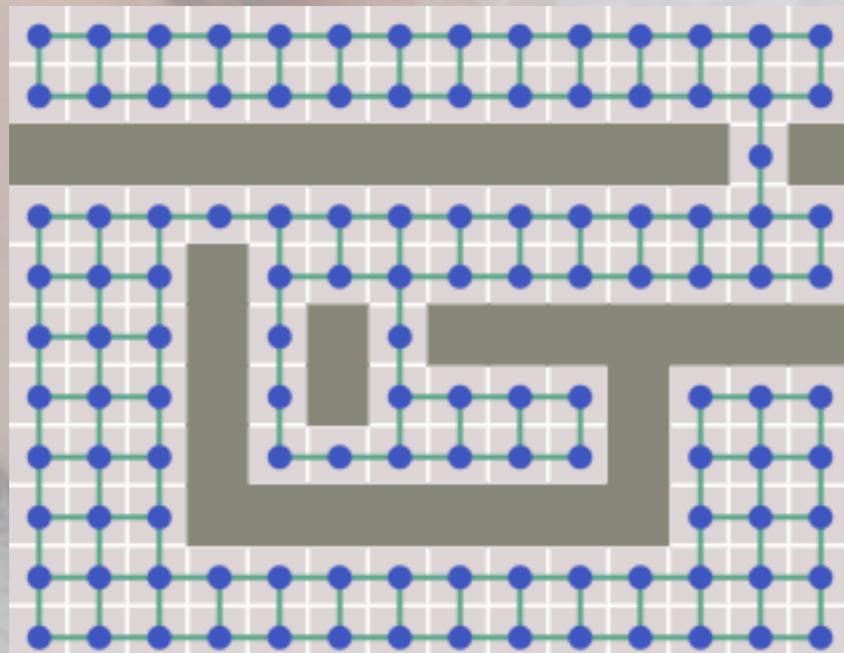
A* Search

- Like Dijkstra's but prefers neighbours that get closer to goal

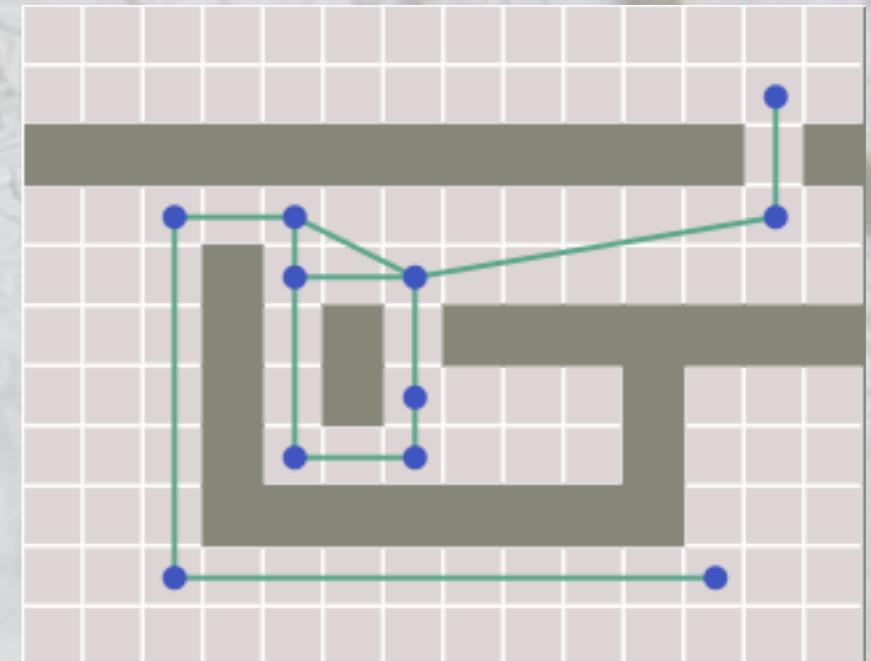


Map Representation

- Map is a graph of connected nodes (the neighbours)
- Grid is converted into graph structure prior to running search



Waypoints would be more efficient...



Implementation

- Keep a queue of “frontier” nodes (unexplored neighbours)
- For each frontier node in turn, explore
- Explore = if goal then done else add all neighbours to frontier
- Keep a “came-from” map of current-node → previous-node
- “came-from” map used to reconstruct path at the end

This is essentially BFS! Dijkstra & A* simply modify the order in which the frontier nodes are explored

Challenges in Clojure

- Clojure likes things to be immutable
- Frontier queue is modified as its being consumed
- Cannot simply map/reduce over a frontier node sequence

Solutions

- Use loop/recur
 - Manually loop through frontier, allowing us to add extra nodes when we recur, if needed
- Use iterate
 - Allows us to generate a new state (frontier queue etc) given the previous state

$f(\text{previous}) = \text{successor}$

Loop/recur

```
(loop [frontier [start]
       came-from {start :start}]
  (let [current (first frontier)]
    (if (not= current goal)
        (let [neighbours (get-neighbours current)]
          (recur (concat (next frontier) neighbours)
                 (into came-from
                       (for [neighbour neighbours]
                         [neighbour current])))))
        came-from)))
```

get-neighbours

- Gets the current node from the global graph
- Gets all neighbour node keys
- Removes neighbours already contained in comes-from map

iterate

```
(take 5 (iterate (fn [x] (* x 2)) 1))
```

((fn [x] (* x 2)) 1))	=> 1
((fn [x] (* x 2)) 2))	=> 2
((fn [x] (* x 2)) 4))	=> 4
((fn [x] (* x 2)) 8))	=> 8
	=> 16

=

```
(1 2 4 8 16)
```

iterate

```
(-> { :frontier [start]
      :came-from {start :start}}
  (iterate
    (fn [{:keys [frontier came-from]}]
      (let [current (first frontier)]
        (when (not= current goal)
          (let [neighbours (get-neighbours current)]
            {:frontier (concat (next frontier)
                               neighbours)
             :came-from (into came-from
                               (for [neighbour neighbours]
                                 [neighbour current]))})))
    (take-while identity)
    last
    :came-from))
```

Dijkstra's Algorithm

- Takes “cost” of a neighbour into account and explores lowest cost first

```
(loop [frontier (priority-queue start)
      came-from {start :start}
      cost-so-far {start 0}]
  (let [current (first (peek frontier))]
    (if (not= current goal)
        (let [neighbours (get-neighbours current)
              costs (get-costs current neighbours
                                 cost-so-far)]
          (recur (into (pop frontier)
                      (for [[node-id cost] costs]
                        [node-id cost])))
                ...
                (into cost-so-far costs)))
```

get-costs

```
(defn get-costs
  [current neighbours cost-so-far]
  (->
    (for [neighbour neighbours]
      (let [new-cost (+ (get cost-so-far current)
                         (get-cost current neighbour))]
        (when (or (not (contains? cost-so-far
                                    neighbour))
                  (< new-cost (get cost-so-far
                                    neighbour))))
          [neighbour new-cost]))))))
```

A* Search

- One-line change to Dijkstra's algorithm! (...plus a heuristic fn)

```
...
(recur
  (into (pop frontier)
    (for [[node-id cost] costs]
      [node-id (+ cost (heuristic goal node-id))]))
...
(defn heuristic
  "Manhattan distance on a square grid"
  [goal current]
  (let [[ax ay] goal
        [bx by] current]
    (+ (Math/abs (- ax bx))
       (Math/abs (- ay by))))))
```

Why does it work?

- Dijkstra calculates the total cost from start to neighbours
- Frontier queue is a priority queue
- Calculated cost used as priority in frontier queue
- Lowest cost items are taken from priority queue first
- A* adds the heuristic (in this case, manhattan distance) to the cost, forcing the queue to prefer “closer” neighbours

Resources

- Amit Patel/Red Blob Games

<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

- Sample code for this talk (uses “iterate” not “loop/recur”)

<https://github.com/danielytics/pathfinding-talk>

Questions?

